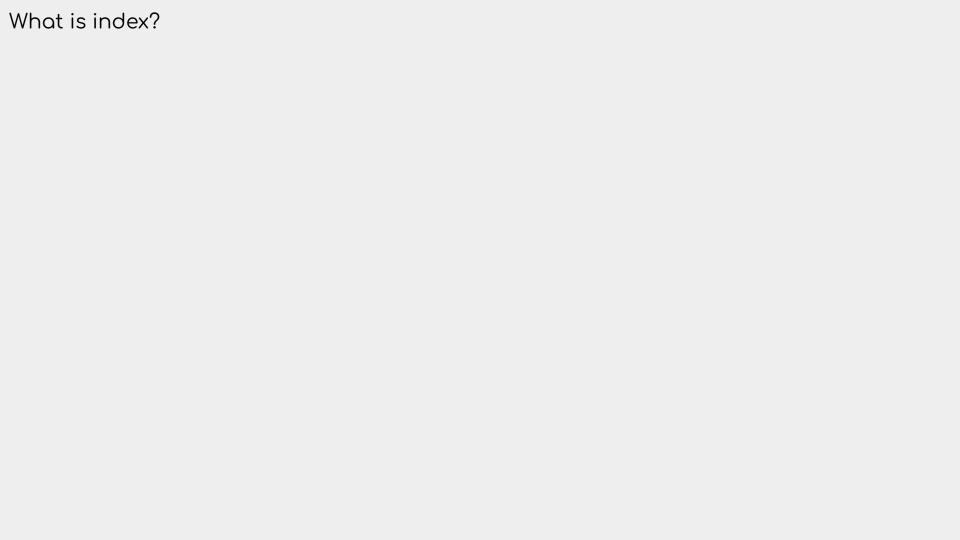
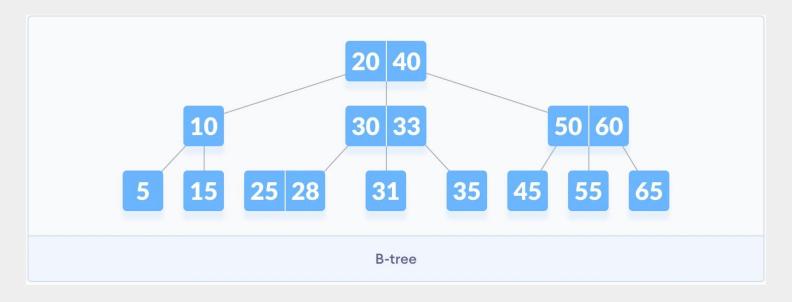
# MongoDB Index 101

2022-06-29





## Data Structure



B-tree is a special type of **self-balancing search tree** in which each node can contain more than one key and can have more than two children.

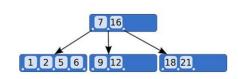
# B-Tree vs Binary Tree

#### **B-tree and Binary Tree Data Structure Comparison**

Both structures operate in the average in  $O(\log n)$  time. Note that in the worst case, B-tree, at  $O(\log n)$ , is faster than Binary Search Tree at O(n).

# B-tree and Binary Search Tree Data Structures

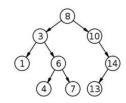
B-tree



	Average	Worst Case
Space	O(n)	O(n)
Search	O(log n)	O(log n)
Insert	O(log n)	O(log n)
Delete	O(log n)	O(log n)

B-tree is optimized for systems that read and write large blocks of data. It is commonly used in databases and file systems.

## **Binary Search Tree**



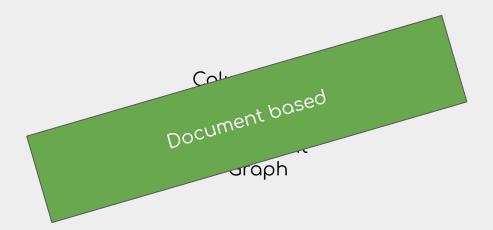
	Average	Worst Case
Space	O(n)	O(n)
Search	O(log n)	O(n)
Insert	O(log n)	O(n)
Delete	O(log n)	O(n)

The major advantage of binary search trees over other data structures is that the related sorting algorithms and search algorithms such as in-order traversal can be very efficient.

MongoDB Database Type

Columns, Row Columnar Key-Value Document Graph

# MongoDB Database Type



## Demo Collection

```
"id" : int 1000

"first_name" : string "Geneva"

"last_name" : string "Kopacek"

"email" : string "gkopacekrr@squarespace.com"

"gender" : string "Agender"

"language" : string "Montenegrin"

"is_active" : bool false
}
```

Database: test Collection: user Documents: 1,000k

## Default Index

```
"_id" : ObjectId("62bbdd5a6978e39c78e3bf87"),

"id" : 1,

"first_name" : "Nikolas",

"last_name" : "Godleman",

"email" : "ngodleman0@webeden.co.uk",

"gender" : "Male",

"language" : "Fijian",

"is_active" : false
}
```

#### getCollectionIndexInfo("users"," id ");

(1)	Key \$	Value
Compression	<b>△</b> ○ (1)	{ name : "_id_" } (6 fields)
	W V	2
	▶ 🖸 key	{ _id : 1 }
■ □ usage stats         Array[1]           ■ □ 0         {name : "_id_", host : "4a334d3d6114:27017" } (5 fields)           □ name         _id_           □ key         {_id : 1 }           □ host         4a334d3d6114:27017           □ accesses         {ops : 2, since : ISODate("2022-06-29T11:54:35.922+07:00") }           □ spec         {v : 2, key : {_id : 1 }, name : "_id_" }           ■ □ index details         {type : "file" } (15 fields)           □ metadata         {formatVersion : 8 }           □ creationString         access_pattern_hint=none, allocation_size=4KB,app_metadata=(file)           □ type         file           □ uri         statistics:table:index-3-4937500220967527413           □ LSM         {12 fields}           □ □ block-manager         {10 fields}           □ □ cache         {22 fields}           □ □ cache         {80 fields}           □ □ cache_walk         {21 fields}           □ □ cache_walk         {21 fields}           □ □ compression         {13 fields}           □ □ compression         {34 fields}           □ □ concolilation         {36 fields}           □ □ ceconciliation         {36 fields}	"" name	_id_
Image	indexSize	65,536 (64.0KB)
iname	■ usage stats	Array[1]
Cache_walk   Cac	<b>⊿</b> □ 0	{ name : "_id_", host : "4a334d3d6114:27017" } (5 fields)
### host	"" name	_id_
Cache   Cach	▶ 🖸 key	{_id:1}
Cache	host	4a334d3d6114:27017
Image:   The composition   T	▶ ☼ accesses	{ ops : 2, since : ISODate("2022-06-29T11:54:35.922+07:00") }
■	▶ 🖸 spec	{ v : 2, key : { _id : 1 }, name : "_id_" }
	■ index details	{ type : "file" } (15 fields)
	▶  ☐ metadata	{ formatVersion : 8 }
statistics:table:index-34937500220967527413	creationString	access_pattern_hint=none,allocation_size=4KB,app_metadata=(f
C   LSM   {12 fields}	type type	file
▶ ☑ block-manager       {10 fields}         ▶ ☑ cache       {22 fields}         ▶ ☑ cache _ walk       {60 fields}         ▶ ☑ cache_walk       {21 fields}         ▶ ☑ checkpoint-cleanup       {"pages added for eviction" : 0, "pages removed" : 0, "pages skipled to the compression         ★ ☑ cursor       {34 fields}         ▶ ☑ reconciliation       {36 fields}         ▶ ☑ session       {} (4 fields)	"" uri	statistics:table:index-34937500220967527413
S	D LSM	{12 fields}
■	▶ □ block-manager	{10 fields}
D C cache_walk         {21 fields}           D C cache_wolnt-cleanup         {"pages added for eviction" : 0, "pages removed" : 0, "pages skipp           D C compression         {13 fields}           D C cursor         {34 fields}           D C reconciliation         {36 fields}           D C session         {} {4 fields}	btree	{22 fields}
↓ Cl checkpoint-cleanup       { "pages added for eviction" : 0, "pages removed" : 0, "pages skipt         ↓ Cl compression       {13 fields}         ↓ Cl cursor       {34 fields}         ↓ Cl reconciliation       {36 fields}         ↓ Cl session       {} (4 fields)	□ cache	{60 fields}
↓ © compression       {13 fields}         ↓ © cursor       {34 fields}         ↓ © reconciliation       {36 fields}         ↓ © session       {} (4 fields)		{21 fields}
↓ CD cursor       {34 fields}         ↓ CD reconciliation       (36 fields)         ↓ CD session       {} (4 fields)	▶ ☼ checkpoint-cleanup	{ "pages added for eviction" : 0, "pages removed" : 0, "pages skipp
	▶ ☐ compression	{13 fields}
	□ cursor	{34 fields}
	▶ ☐ reconciliation	{36 fields}
		{ } (4 fields)
	▶ □ transaction	{13 fields}

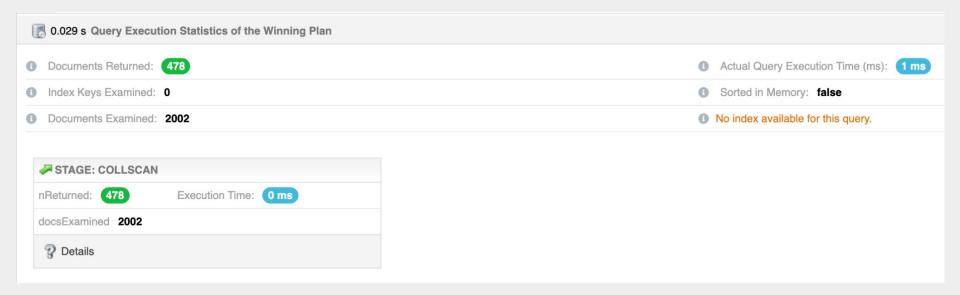
# Query with no Index

```
db.users.find({"is active": false})
```

```
# slow query
# expensive operation
# long command operation
# collection scan
```

# Explain();

db.users.find({"is active": true}).explain(true)



# Simple Index explanation

#### Consider an index { a:1} Just only "a" field in index

```
{ __id: ObjectId(), a: 1, b: "ab" }
{ __id: ObjectId(), a: 1, b: "cd" }
{ __id: ObjectId(), a: 1, b: "ef" }
{ __id: ObjectId(), a: 2, b: "jk" }
{ __id: ObjectId(), a: 2, b: "lm" }
{ __id: ObjectId(), a: 2, b: "no" }
{ __id: ObjectId(), a: 3, b: "pq" }
{ __id: ObjectId(), a: 3, b: "rs" }
{ __id: ObjectId(), a: 3, b: "tv" }
```

If you query for { a: 2, b: "no" } MongoDB must scan 3 documents in the collection to return the one matching result.

# Design the Index

```
# Know your application (read/write heavy)
# Know you query
```

## Example

An API for order services

- 1. Find by order\_id
- 2. Find by order status
- 3. Find by user\_id
- 4. ..

#### **Production Practices**

## **High Cardinality**

The index cardinality refers to how many possible values there are for a field. The field sex only has two possible values. It has a very low cardinality. Other fields such as names, usernames, phone numbers, emails, etc. will have a more unique value for every document in the collection, which is considered high cardinality.

#### Compound Index

Try to create indexes on high-cardinality keys or put high-cardinality keys first in the compound index.

## Selectivity

```
{_id:ObjectId(), name:"John", sex:"male"}
{_id:ObjectId(), name:"Rich", sex:"male"}
{_id:ObjectId(), name:"Mose", sex:"male"}
{_id:ObjectId(), name:"Sami", sex:"male"}
{_id:ObjectId(), name:"Cari", sex:"female"}
{_id:ObjectId(), name:"Mary", sex:"female"}
```

#### Case#1

If your index is {name:1}, If you run the query name: "John", sex: "male"}

You will have to scan 1 document. Because you allowed MongoDB to be selective.

#### Case#2

If your index is {sex:1}, If you run the query

{sex: "male", name: "John"}

You will have to scan 4 documents.

#### Take a Break with Book

#### **Index Cardinality**

Cardinality refers to how many distinct values there are for a field in a collection. Some fields, such as "gender" or "newsletter opt-out", might only have two possible values, which is considered a very low cardinality. Others, such as "username" or "email", might have a unique value for every document in the collection, which is high cardinality. Still others fall somewhere in between, such as "age" or "zip code".

In general, the greater the cardinality of a field, the more helpful an index on that field can be. This is because the index can quickly narrow the search space to a much smaller result set. For a low cardinality field, an index generally cannot eliminate as many possible matches.

For example, suppose we had an index on "gender" and were looking for women named Susan. We could only narrow down the result space by approximately 50% before referring to individual documents to look up "name". Conversely, if we indexed by "name", we could immediately narrow down our result set to the tiny fraction of users named Susan and then we could refer to those documents to check the gender.

As a rule of thumb, try to create indexes on high-cardinality keys or at least put highcardinality keys first in compound indexes (before low-cardinality keys).

#### Using explain() and hint()

As you have seen above, explain() gives you lots of information about your queries. It is one of the most important diagnostic tools there is for slow queries. You can find out which indexes are being used and how by looking at a query's explain. For any query, you can add a call to explain() at the end (the way you would add a sort() or lim it(), but explain() must be the last call).

There are two types of explain() output that you'll see most commonly: indexed and non-indexed queries. Special index types may create slightly different query plans, but most fields should be similar. Also, sharding returns a conglomerate of explain()s (as covered in Chapter 13), as it runs the query on multiple servers.

The most basic type of explain() is on a query that doesn't use an index. You can tell that a query doesn't use an index because it uses a "BasicCursor". Conversely, most queries that use an index use a "BtreeCursor" (some special types of indexes, such as geospatial indexes, use their own type of cursor).

The output to an explain() on a query that uses an index varies, but in the simplest case, it looks something like this:

```
> db.users.find({"age" : 42}).explain()
    "cursor" : "BtreeCursor age_1_username_1",
   "isMultiKey" : false,
```

Powerful and Scalable Data Storage The Definitive Guide O'REILLY® Kristina Chodorow

## Planner & Stage

#### Planner

```
queryPlanner = winning plan
rejected plans = mongo optimizer reject the plan
executionStats = show execution stats
allPlansExecution = show all plan
```

## Stage

- COLLSCAN for a collection scan
- IXSCAN for scanning index keys
- FETCH for retrieving documents
- SHARD\_MERGE for merging results from shards
- SHARDING\_FILTER for filtering out orphan documents from shards

#### Usecase#1

```
" id" : <a href="mailto:objectId">ObjectId</a>("62bc2eca6978e39c78e3c757"),
"id" : 1000,
"first name" : "Geneva",
"last name" : "Kopacek",
"email" : "gkopacekrr@squarespace.com",
"gender" : "Agender",
"language" : "Montenegrin",
"is active" : false
" id" : <a href="mailto:objectId">objectId</a>("62bc2eca6978e39c78e3c756"),
"id" : 999,
"first name" : "Samaria",
"last name" : "Tumility",
"email" : "stumilityrg@indiegogo.com",
"gender" : "Female",
"language" : "Kannada",
"is active" : true
```

```
db.getSiblingDB("demo")
 db.users.createIndex({language: 1, is active: 1});
 1 use "demo"
 3 db.users.find({"language": "Japanese"}).explain(true)
3 0.066 s Query Execution Statistics of the Winning Plan

    Actual Query Execution Time (ms): 11 ms

    Documents Returned: 14

Index Keys Examined: 14
                                          Sorted in Memory: false
Ocuments Examined: 14
                                          Execution Time: 7 ms 100%
 nReturned: 14
  Details
  STAGE: IXSCAN
 nReturned: 14
                  Execution Time: 0 ms
 indexName REGULAR (1) language_1_is_active_1
  Details
```

## Usecase#2

```
" id" : <a href="mailto:objectId">ObjectId</a>("62bc2eca6978e39c78e3c757"),
"id" : 1000,
"first name" : "Geneva",
"last name" : "Kopacek",
"email" : "gkopacekrr@squarespace.com",
"gender" : "Agender",
"language" : "Montenegrin",
"is active" : false
" id" : <a href="mailto:objectId">objectId</a>("62bc2eca6978e39c78e3c756"),
"id" : 999,
"first name" : "Samaria",
"last name" : "Tumility",
"email" : "stumilityrg@indiegogo.com",
"gender" : "Female",
"language" : "Kannada",
"is active" : true
```

```
db.getSiblingDB("demo")
db.users.createIndex({language: 1, is active: 1});
 1 use "demo"
 3 db.users.find({"language": "Japanese", "is_active": true}).explain(true)
R 0.035 s Query Execution Statistics of the Winning Plan

    Documents Returned: 6

    Actual Query Execution Time (ms): 2 ms

Index Keys Examined: 6
                                               Sorted in Memory: false

    Documents Examined: 6

                                               Query used the following index: REGULAR (1) language_1_is_active_1
  Execution Time: 0 ms
 nReturned: 6
  Details
  Execution Time: 0 ms
 nReturned: 6
 indexName REGULAR (1) language 1 is active 1
  Details
```

## Usecase#3

```
" id" : <a href="mailto:objectId">ObjectId</a>("62bc2eca6978e39c78e3c757"),
"id" : 1000,
"first name" : "Geneva",
"last name" : "Kopacek",
"email" : "gkopacekrr@squarespace.com",
"gender" : "Agender",
"language" : "Montenegrin",
"is active" : false
" id" : <a href="mailto:objectId">ObjectId</a>("62bc2eca6978e39c78e3c756"),
"id" : 999,
"first name" : "Samaria",
"last name" : "Tumility",
"email" : "stumilityrg@indiegogo.com",
"gender" : "Female",
"language" : "Kannada",
"is active" : true
```

```
db.getSiblingDB("demo")
db.users.createIndex({language: 1, is active: 1});
     use "demo"
 3 db.users.find({|"is_active": true}).explain(true)
7 0.032 s Query Execution Statistics of the Winning Plan

    Documents Returned: 478

    Actual Query Execution Time (ms): 1 ms

  Index Keys Examined: 0
                                                    Sorted in Memory: false
Documents Examined: 1001
                                                    No index available for this query.
  ₹ STAGE: COLLSCAN
 nReturned: 478
                       Execution Time: 0 ms
  docsExamined 1001
  Details
```